



**Neural networks through ice cream
sales**

Objectives:

- To *fluently* use the *vocabulary* of neural networks
- To connect *familiar past algorithms* to *neural networks*, and therefore demystify the math
- To practice *drawing* and *specifying* the parameters of a neural net
- To list the *parameters* one can *adjust* when building a neural net



Scenario

You own a chain of ice cream stores.

You want to build a model that will predict the sales numbers of a store, given the store's location, pricing of product, and perceived quality of the product.

Simpler models haven't produced great results, so you want to try a neural network. Plus, neural networks sound fancy. You like fancy.






Problem summary




Variables


Location



Pricing




Perceived Quality




Variables


Location



Pricing




Perceived Quality




Target

Sales




Variables


Location



Pricing




Perceived Quality




Target

Sales




Variables


Location



Pricing




Perceived Quality



Predict


Target

Sales




Variables


Location



Pricing



Perceived Quality




Predict

**Using a Neural
Network**

Target

Sales






Vocabulary




Using a Neural Network

Variables


Location



Pricing




Perceived Quality



Target


Sales




Using a Neural Network

Variables


Location



Pricing




Perceived Quality



input
layer

Target


Sales




Using a Neural Network

Variables


Location



Pricing




Perceived Quality



**input
layer**
nodes: 3

Target


Sales




Using a Neural Network

Variables


Location



Pricing




Perceived Quality



**input
layer**
nodes: 3

Target

Sales




**output
layer**


Using a Neural Network

Variables


Location



Pricing



Perceived Quality




**input
layer**

nodes: 3

Target

Sales




**output
layer**


nodes: 1

Variables


Location



Pricing







Perceived Quality



**input
layer**


nodes: 3

Using a Neural Network



Target

Sales




**output
layer**


nodes: 1

Variables


Location



Pricing







Perceived Quality



**input
layer**

nodes: 3


Using a Neural Network



**hidden
layer**

Target

Sales




**output
layer**


nodes: 1

Variables


Location



Pricing







Perceived Quality



**input
layer**

nodes: 3

Using a Neural Network




**hidden
layer**

nodes: 4

Target

Sales



**output
layer**

node: 1



How many layers in our neural network?



Keep this in mind as we build neural networks




Draw this out:

We want to build a neural network using gender(assume binary), years of education, marital status(single vs wed), and years of employment to predict income.


We are going to use two hidden layers. The first one will have three nodes and the second will have 5.

Draw and compare w neighbors.


The math behind networks is not that scary




We need some notation to make this work



May have seen diagrams like this




For simplicity we are only going to focus on one layer



input
layer
nodes: 3

hidden
layer
nodes: 4


And specifically the first node in the hidden layer



input
layer
nodes: 3

hidden
layer
nodes: 4


Problem: this common diagram isn't representative



input
layer
nodes: 3

hidden
layer


What is shown as one is really three

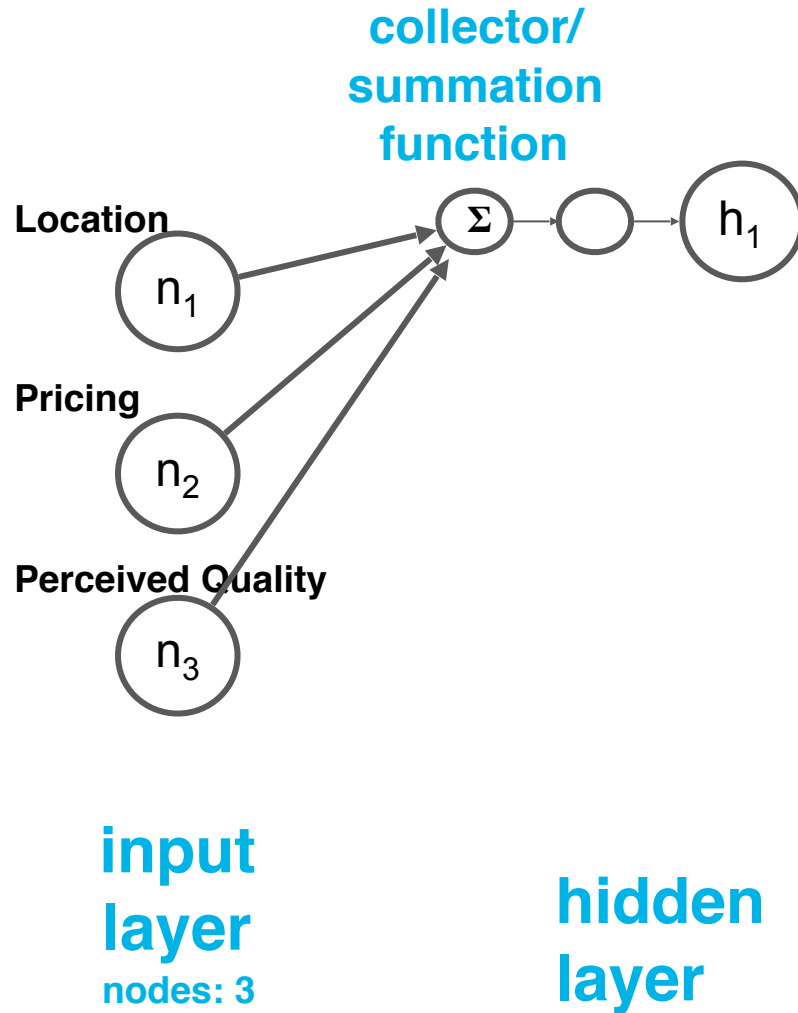


input
layer
nodes: 3


hidden
layer

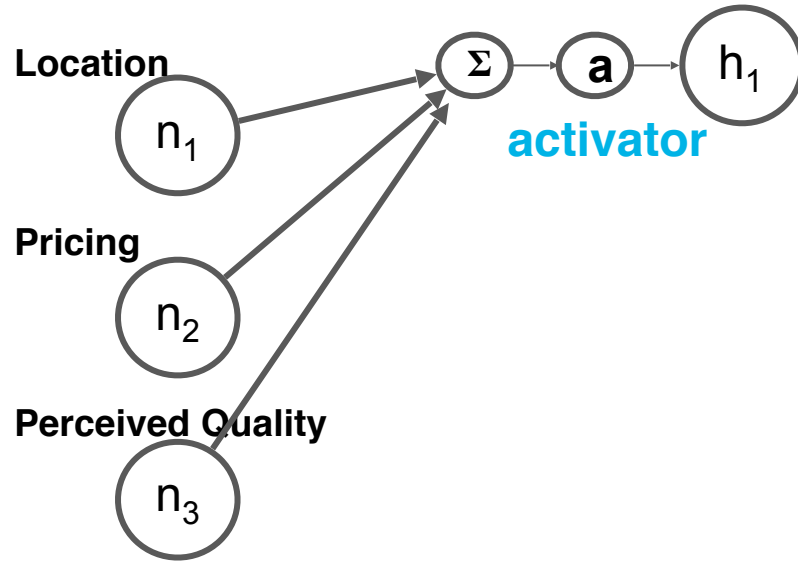
Each node that receives input from previous nodes actually has **three** parts





Each node that receives input from previous nodes actually has **three** parts




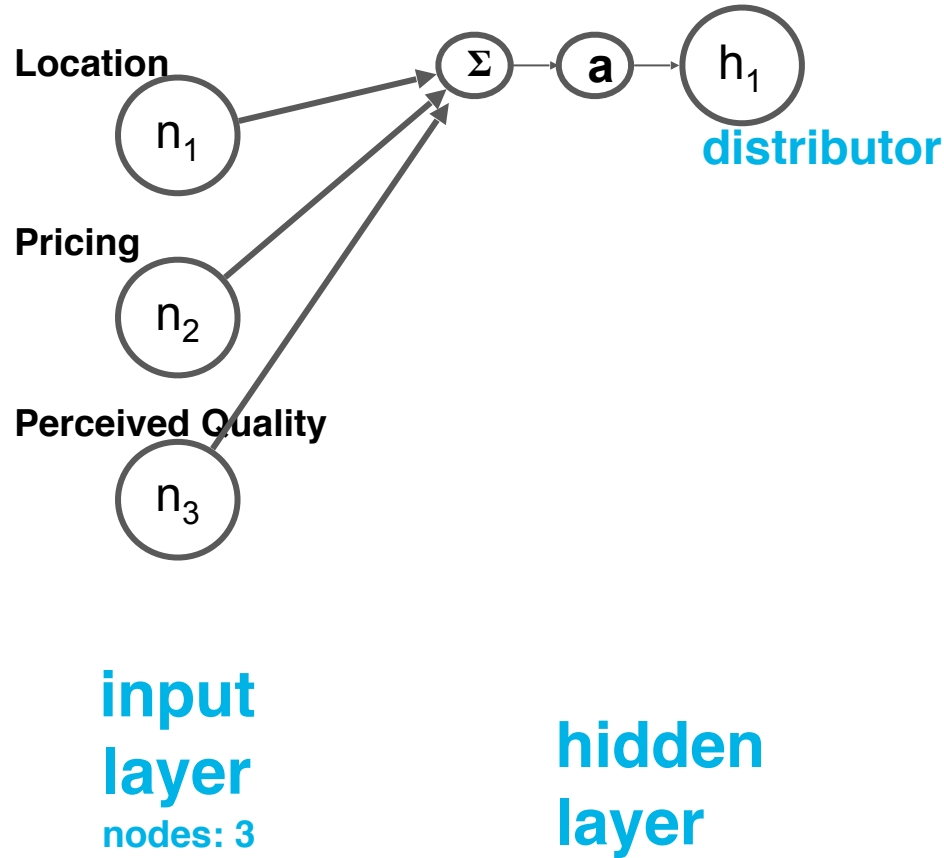


input
layer
nodes: 3


hidden
layer

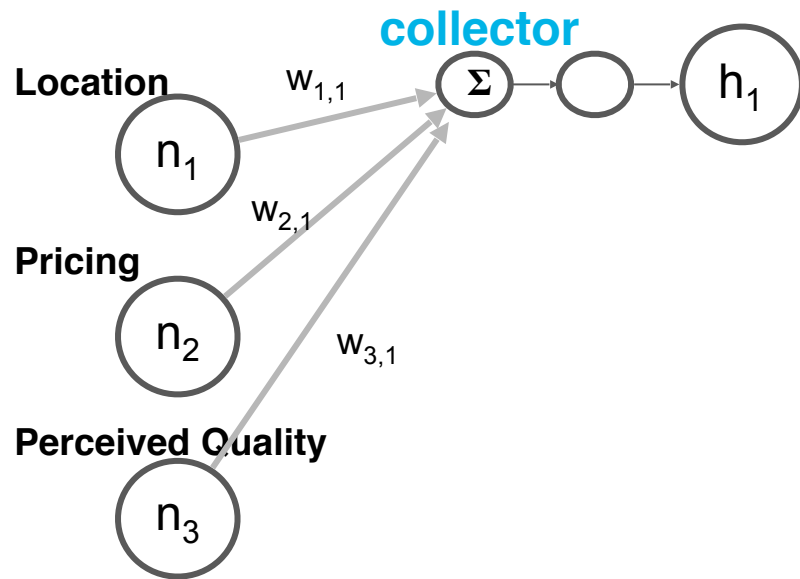
Each node that receives input from previous nodes actually has **three** parts





Each node that receives input from previous nodes actually has **three** parts

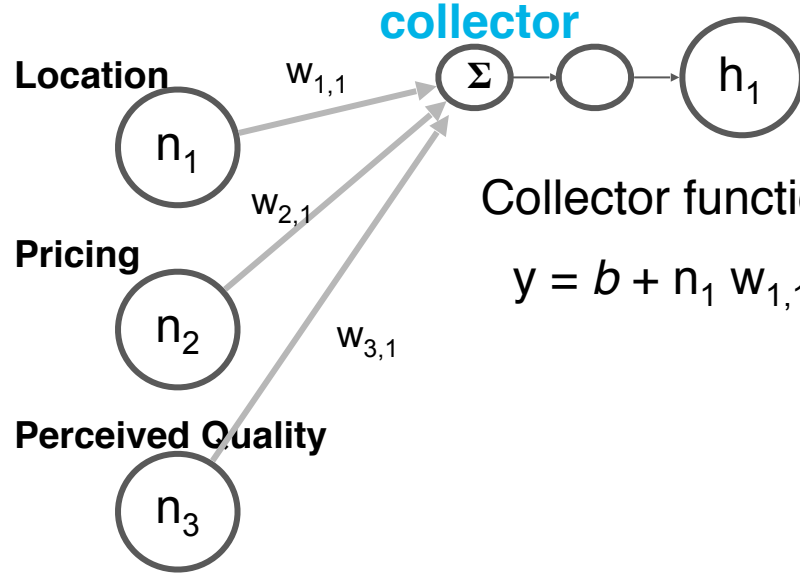




input
layer
nodes: 3

hidden
layer

This function should look familiar

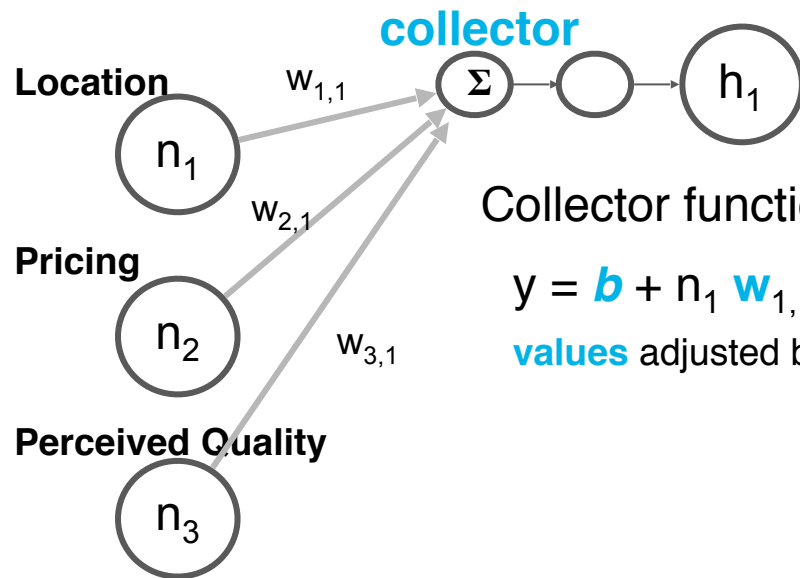


Collector function:

$$y = b + n_1 w_{1,1} + n_2 w_{2,1} + n_3 w_{3,1}$$

**input
layer**
nodes: 3

**hidden
layer**



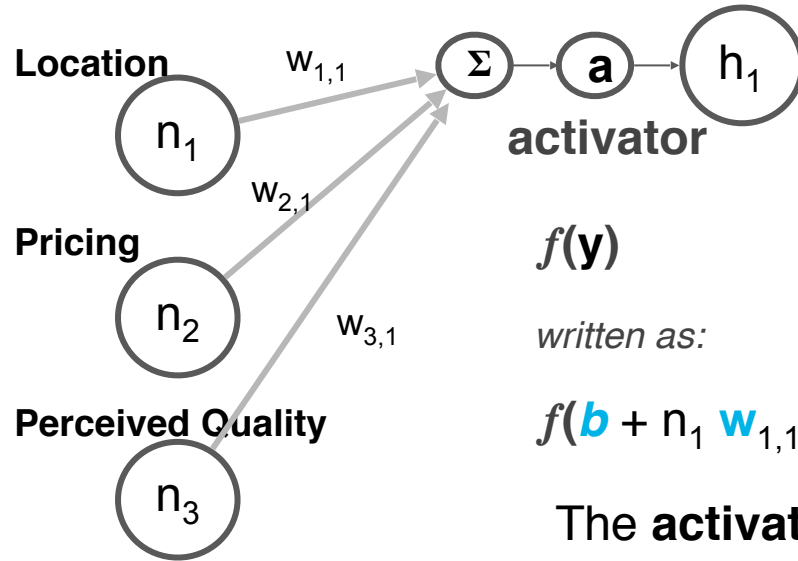
Collector function:

$$y = b + n_1 w_{1,1} + n_2 w_{2,1} + n_3 w_{3,1}$$

values adjusted by the algorithm

**input
layer**
nodes: 3

**hidden
layer**



The **activator** is a function chosen by **you** that takes the output of the collector as input.

$f(y)$

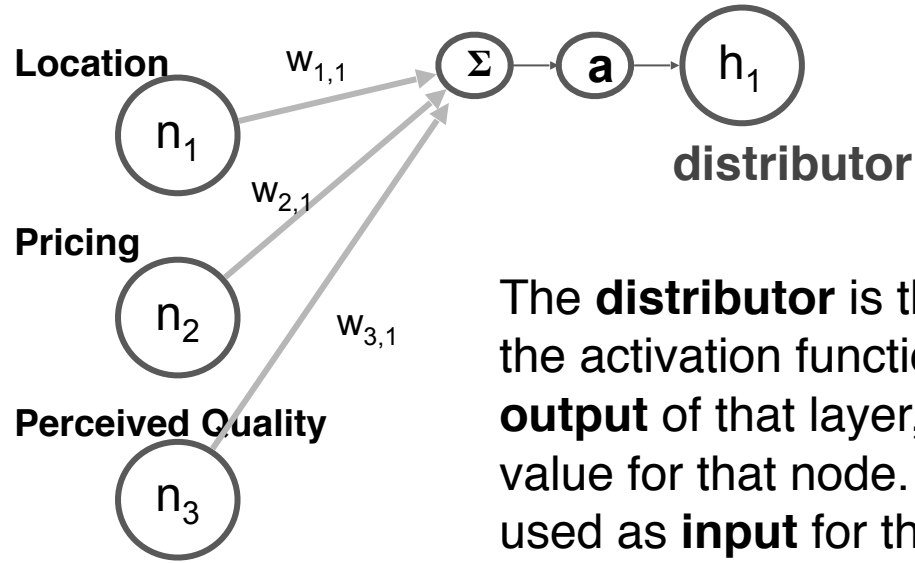
written as:

$$f(b + n_1 w_{1,1} + n_2 w_{2,1} + n_3 w_{3,1})$$

The **activator** is specified for **each** layer. Nodes within a layer all use the same activation function

**input
layer**
nodes: 3

**hidden
layer**

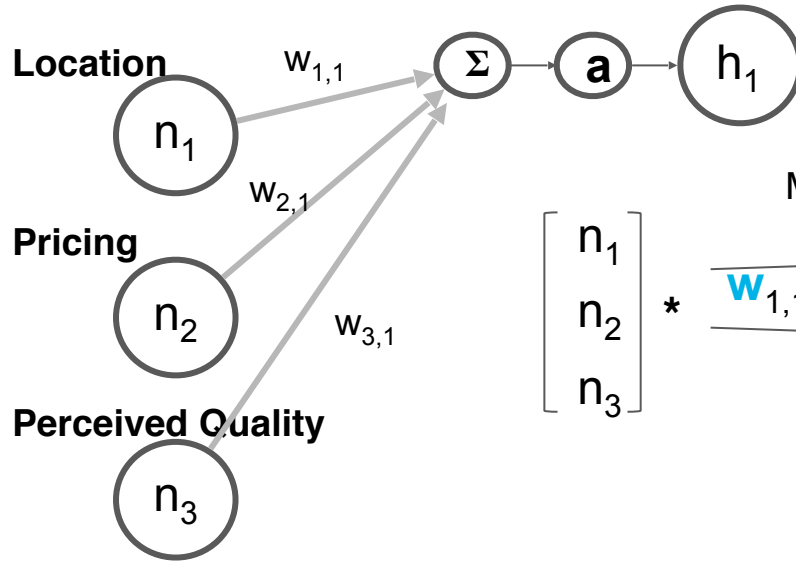


The **distributor** is the **output** of the activation function. It is the **output** of that layer, the final value for that node. It is then used as **input** for the next layer.

input
layer
nodes: 3

hidden
layer

For the math lovers in the room:




Matrix notation for one hidden node

$$\begin{bmatrix} n_1 \\ n_2 \\ n_3 \end{bmatrix} * \overline{\mathbf{w}_{1,1} \mathbf{w}_{2,1} \mathbf{w}_{3,1}} + \mathbf{b}_1 = z_1 \rightarrow f_{a1}(y_1) \rightarrow h_1$$

input
layer
nodes: 3

hidden
layer

The full math for a layer:



**input
layer**
nodes: 3

**hidden
layer**
nodes: 4

Matrix notation for one hidden layer

$$\begin{bmatrix} n_1 \\ n_2 \\ n_3 \end{bmatrix} * \begin{bmatrix} w_{1,1} & w_{2,1} & w_{3,1} \\ w_{1,2} & w_{2,2} & w_{3,2} \\ w_{1,3} & w_{2,3} & w_{3,3} \\ w_{1,4} & w_{2,4} & w_{3,4} \end{bmatrix}^T + \begin{bmatrix} b_1 \\ b_2 \\ b_3 \\ b_4 \end{bmatrix} = \begin{bmatrix} y_1 \\ y_2 \\ y_3 \\ y_4 \end{bmatrix}$$

$$\begin{bmatrix} f(y_1) \\ f(y_2) \\ f(y_3) \\ f(y_4) \end{bmatrix} = \begin{bmatrix} h_1 \\ h_2 \\ h_3 \\ h_4 \end{bmatrix}$$

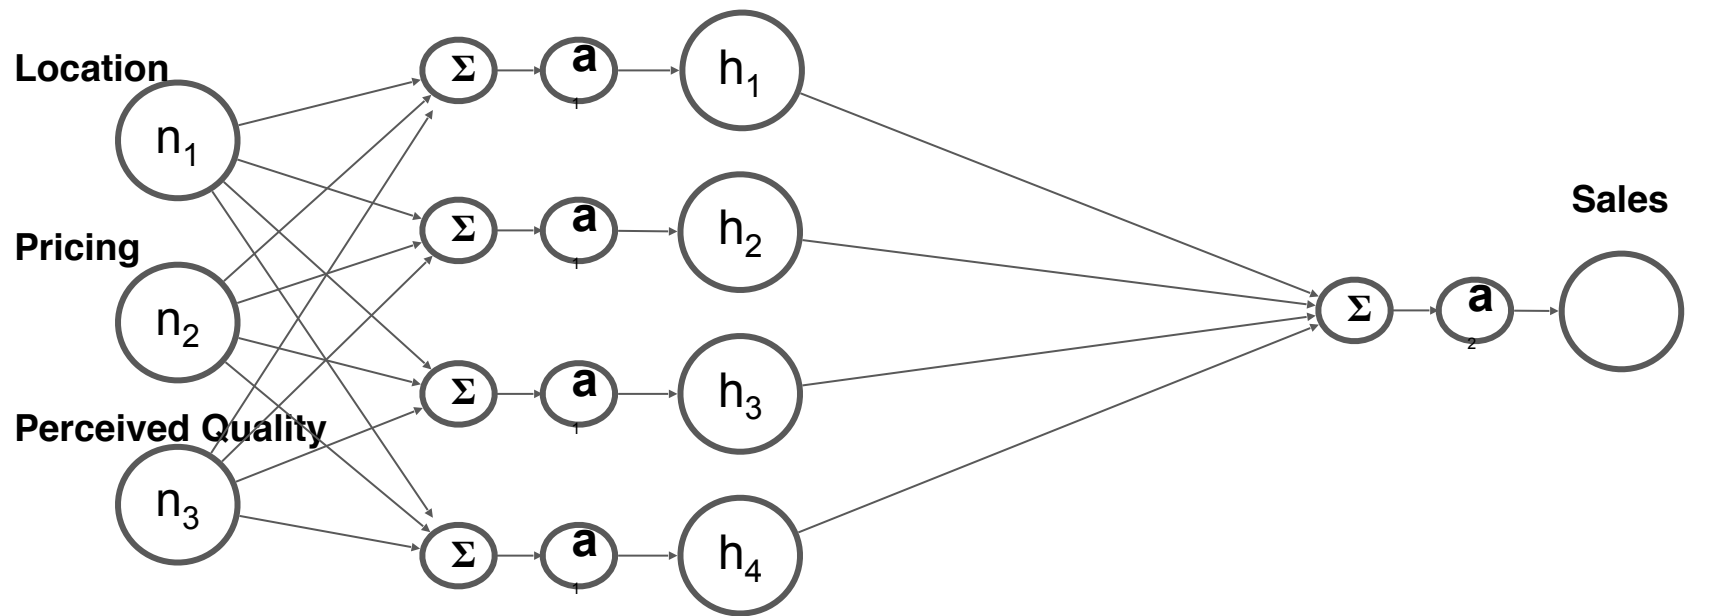


When creating a layer, what are the specs that **you** choose?



Summary

Hidden layer variables	You define	Computer figures out
		✓
weights	✓	
activation function		✓
bias	✓	
number of nodes		



input
layer
nodes: 3

hidden
layer
nodes: 4

output
layer
one node

Summary so far:


What you choose when building a neural network:

At the **network** level:

- Number of input variables
- Number of hidden layers
- If it is a classification or regression problem

At the **layer** level:

- The number of nodes
- The activation function





What else can we adjust?



What you choose when building a neural network:

At the **network** level:

- Number of input variables
- Number of hidden layers
- If it is a classification or regression problem
- **Batch size**
- **Number of epochs**
- **Learning rate & optimizer**
- **Regularization type and lambda**

At the **layer** level:

- The number of nodes
- The activation function



Batches and Epochs are about data processing

That's a lot of math and a lot of data.
The dataset is split into chunks and passed through the network one chunk at a time.

Batch defines the number of observations in each “chunk”.


Epochs are how many times you want the whole dataset to go through the network.

All the **batches** = one **epoch**.




Learning rate is from gradient descent

Scenario	Cost/Loss Function
Regression	MSE
Binary classification	Cross-Entropy (Logarithmic loss)
Multi-class classification	Softmax of Cross-Entropy




FP




Optimizer is how the gradient is calculated

Options
sgd
rmsprop
Adagrad
Adadelata
Adam
And more!



FP



Regularization - adjusts the weights by layer

Adding a set penalization term at each collector node.

You can choose no regularization, lasso, or ridge.

Write in sentences what this code does

```
model = Sequential()
model.add(Dense(32, activation='relu', input_dim=100))
model.add(Dense(1, activation='sigmoid'))
model.compile(optimizer='rmsprop',
              loss='binary_crossentropy',
              metrics=['accuracy'])


# Generate dummy data
import numpy as np
data = np.random.random((1000, 100))
labels = np.random.randint(2, size=(1000, 1))

# Train the model, iterating on the data in batches of 32 samples
model.fit(data, labels, epochs=10, batch_size=32)
```

Even without learning Keras explicitly, you should be able to recognize keywords and concepts based on this review.

Dense = fully connected to all previous nodes

While we
have
covered
this:



Know this
is a whole
additional
area.

